

19) Convolutional Neural Networks: Padding, Strides, Max-Pooling

Vitor Kamada


August 2019

Chollet (2018): Ch 5.1

<https://www.manning.com/books/deep-learning-with-python>

<https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/5.1-introduction-to-convnets.ipynb>

Global vs Local Patterns

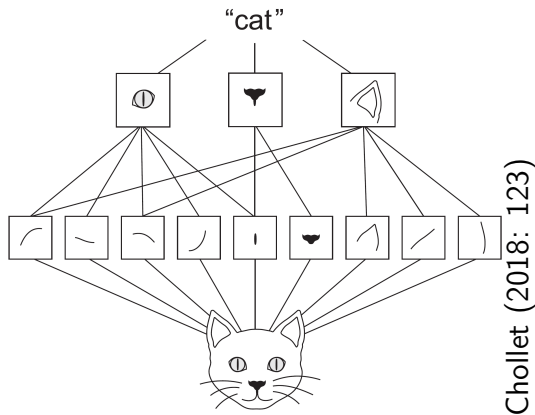


3D tensors (feature maps) = $(28, 28, 1)$

Spatial Axes (Height and Width) = 28×28 pixels

Depth Axis (Channels) = 1 (levels of gray)


Spatial Hierarchies of Patterns



Depth is a parameter of the layer, and no longer stand for specific colors

Filters encode concepts: eyes, face

Feature Map of Size (28, 28, 1)

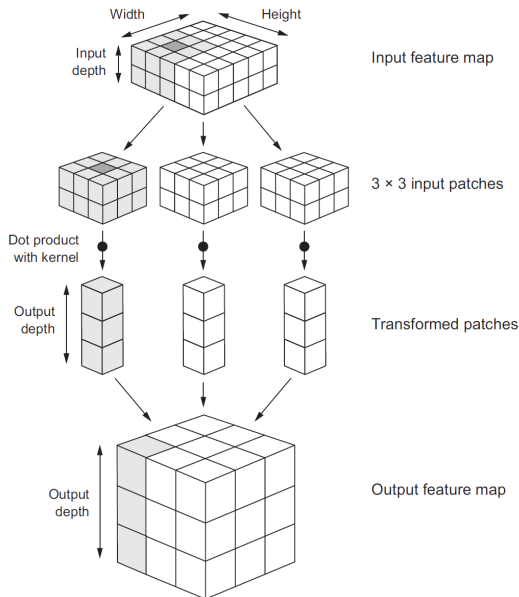


Size of the patches: 3×3

Depth of the output feature map: 32

Outputs a Feature Map of Size (26, 26, 32)



How Convolution Works?



Chollet (2018: 125)

3×3 Patches in a 5×5 Input Feature Map

Border Effect: 5×5→3×3




Chollet (2018: 126)

Padding a 5×5 input to extract 25 3×3 patches

In Conv2D layers:

padding = “valid”: means no padding



Chollet (2018: 126)

3×3 Convolution Patches with 2×2 Strides

Stride: Distance between two Successive Windows

	1		2	
	3		4	

	1	


	2	

	3	

	4	

Chollet (2018: 126)

Max-pooling Operation



Source: Stanford's CS231n github

```
from keras import layers
from keras import models
```

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928

$$(\text{output_channels}) * (\text{input_channels} * \text{window_size} + 1)$$

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
Total params: 55,744		
Trainable params: 55,744		
Non-trainable params: 0		

$$32 * (1 * 3 * 3 + 1) = 320$$

$$64 * (32 * 3 * 3 + 1) = 18496$$

MNIST Digits

```
model.add(layers.Flatten())  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10, activation='softmax'))
```

```
from keras.datasets import mnist  
from keras.utils import to_categorical  
  
(train_images, train_labels), (test_images,  
                                test_labels) = mnist.load_data()  
  
train_images = train_images.reshape((60000, 28, 28, 1))  
train_images = train_images.astype('float32') / 255  
  
test_images = test_images.reshape((10000, 28, 28, 1))  
test_images = test_images.astype('float32') / 255  
  
train_labels = to_categorical(train_labels)  
test_labels = to_categorical(test_labels)
```

```
model.compile(optimizer='rmsprop',  
              loss='categorical_crossentropy', metrics=['accuracy'])  
model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

```
Epoch 1/5  
60000/60000 [=====] - 8s - loss: 0.1766 - acc: 0.9440  
Epoch 2/5  
60000/60000 [=====] - 7s - loss: 0.0462 - acc: 0.9855  
Epoch 3/5  
60000/60000 [=====] - 7s - loss: 0.0322 - acc: 0.9902  
Epoch 4/5  
60000/60000 [=====] - 7s - loss: 0.0241 - acc: 0.9926  
Epoch 5/5  
60000/60000 [=====] - 7s - loss: 0.0187 - acc: 0.9943
```

```
test_loss, test_acc = model.evaluate(test_images, test_labels)
```

```
test_acc
```

0.99129999