

18) Predicting House Prices with Keras

Vitor Kamada

August 2019

Chollet, F. (2018). **Deep Learning with Python**. Ch 3.

<https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/3.7-predicting-house-prices.ipynb>

The Boston Housing Price Dataset

```
import keras
```

```
from keras.datasets import boston_housing  
(train_data, train_targets), (test_data,  
    test_targets) = boston_housing.load_data()
```

```
train_data.shape
```

```
(404, 13)
```

```
test_data.shape
```

```
(102, 13)
```

13 Features

- ① Per capita crime rate.
- ② Proportion of residential land zoned for lots over 25,000 square feet.
- ③ Proportion of non-retail business acres per town.
- ④ Charles River dummy variable (= 1 if tract bounds river; 0 otherwise).
- ⑤ Nitric oxides concentration (parts per 10 million).
- ⑥ Average number of rooms per dwelling.
- ⑦ Proportion of owner-occupied units built prior to 1940.
- ⑧ Weighted distances to five Boston employment centres.
- ⑨ Index of accessibility to radial highways.
- ⑩ Full-value property-tax rate per \$10,000.
- ⑪ Pupil-teacher ratio by town.
- ⑫ $1000 * (B_k - 0.63) ** 2$ where B_k is the proportion of Black people by town.
- ⑬ % lower status of the population.

Normalizing the Data

```
train_targets
```

```
array([[ 15.2,  42.3,  50. ,  
        14.4,  12.1,  17.9,  
        24.1,  27.5,  10.9,  
        34.7,  16.6,  17.5,  
        13.9,  13.1,  20.4,
```

$$Z = \frac{X - \mu}{\sigma}$$

```
mean = train_data.mean(axis=0)  
train_data -= mean  
std = train_data.std(axis=0)  
train_data /= std
```

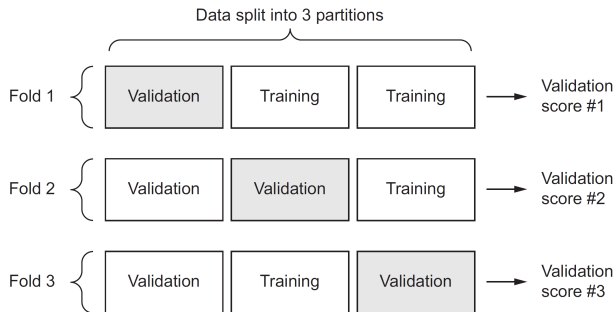
```
test_data -= mean  
test_data /= std
```

Building the Network

```
from keras import models
from keras import layers

def build_model():
    # Because we will need to instantiate
    # the same model multiple times,
    # we use a function to construct it.
    model = models.Sequential()
    model.add(layers.Dense(64, activation='relu',
        input_shape=(train_data.shape[1],)))
    model.add(layers.Dense(64, activation='relu'))
    model.add(layers.Dense(1))
    model.compile(optimizer='rmsprop',
        loss='mse', metrics=['mae'])
    return model
```

K-fold Validation



Chollet (2018:87)

```
import numpy as np
```

```
k = 4
```

```
num_val_samples = len(train_data) // k
```

```
num_epochs = 500
```

```
all_mae_histories = []
```


k=4, epochs = 500

```
for i in range(k):
    print('processing fold #', i)
    # Prepare the validation data: data from partition # k
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
    # Prepare the training data: data from all other partitions
    partial_train_data = np.concatenate(
        [train_data[:i * num_val_samples],
         train_data[(i + 1) * num_val_samples:]],
        axis=0)
    partial_train_targets = np.concatenate(
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]],
        axis=0)
    # Build the Keras model (already compiled)
    model = build_model()
    # Train the model (in silent mode, verbose=0)
    history = model.fit(partial_train_data, partial_train_targets,
                        validation_data=(val_data, val_targets),
                        epochs=num_epochs, batch_size=1, verbose=0)
    mae_history = history.history['val_mean_absolute_error']
    all_mae_histories.append(mae_history)
```


Avoid Overfitting

```
average_mae_history = [  
    np.mean([x[i] for x in all_mae_histories]  
            ) for i in range(num_epochs)]
```

```
import matplotlib.pyplot as plt  
plt.plot(range(1, len(average_mae_history) + 1),  
         average_mae_history)  
plt.xlabel('Epochs').  
plt.ylabel('Validation MAE')  
plt.show()
```



Performance on Test Data

```
# Get a fresh, compiled model.  
model = build_model()  
# Train it on the entirety of the data.  
model.fit(train_data, train_targets,  
          epochs=80, batch_size=16, verbose=0)  
test_mse_score, test_mae_score = model.evaluate(test_data,  
                                                test_targets)
```

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

```
test_mae_score
```

```
2.5532484335057877
```