

17) Sentiment Analysis on Movie Reviews

Vitor Kamada

August 2019

Chollet, F. (2018). **Deep Learning with Python**. Ch 3.

<https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/3.5-classifying-movie-reviews.ipynb>

Internet Movie Database (IMDB)

```
import keras
```

```
from keras.datasets import imdb
```

```
(train_data, train_labels), (test_data,  
    test_labels)= imdb.load_data(num_words=10000)
```

ValueError Traceback (most recent call last)

[<ipython-input-13-00d01dba4bc2>](#) in [<module>\(\)](#)

2

3 (train_data, train_labels), (test_data,

----> 4 test_labels)= imdb.load_data(num_words=10000)

⌵ 2 frames

[/usr/local/lib/python3.6/dist-packages/numpy/lib/format.py](#) in [read_array\(fp, allow_pickle,](#)

694 # The array contained Python objects. We need to unpickle the data.

695 if not allow_pickle:

--> 696 raise ValueError("Object arrays cannot be loaded when "

697 "allow_pickle=False")

698 if pickle_kwargs is None:

ValueError: Object arrays cannot be loaded when allow_pickle=False

SEARCH STACK OVERFLOW

New Code to Load IMDB

```
import numpy as np
# save np.load
np_load_old = np.load

# modify the default parameters of np.load
np.load = lambda *a,**k: np_load_old(*a,
                                     allow_pickle=True, **k)

# call load_data with allow_pickle implicitly set to true
(train_data, train_labels), (test_data,
                             test_labels) = imdb.load_data(num_words=10000)

# restore np.load for future normal usage
np.load = np_load_old
```

Reviews	Total	Positive	Negative
Training	25,000	50%	50%
Testing	25,000	50%	50%

Decode Back to English

```
train_data[0]
```

```
[1,  
14,  
22,  
16,  
43,  
530,
```

```
train_labels[0]
```

```
1
```

```
# word_index is a dictionary mapping words to an integer index  
word_index = imdb.get_word_index()  
# We reverse it, mapping integer indices to words  
reverse_word_index = dict([(value, key) for  
                           (key, value) in word_index.items()])  
# We decode the review; note that our indices were offset by 3  
# because 0, 1 and 2 are reserved indices for "padding",  
# "start of sequence", and "unknown".  
decoded_review = ' '.join([reverse_word_index.get(i - 3, '?')  
                           for i in train_data[0]])
```

```
decoded_review
```

"? this film was just brilliant casting location scenery
story direction everyone's really suited the part they

Vectorize the Data

```
def vectorize_sequences(sequences, dimension=10000):  
    # Create an all-zero matrix of shape  
    # (len(sequences), dimension)  
    results = np.zeros((len(sequences), dimension))  
    for i, sequence in enumerate(sequences):  
        # set specific indices of results[i] to 1s  
        results[i, sequence] = 1.  
    return results
```

```
# Our vectorized training data  
x_train = vectorize_sequences(train_data)  
# Our vectorized test data  
x_test = vectorize_sequences(test_data)
```

```
x_train[0]
```

```
array([ 0.,  1.,  1., ...,  0.,  0.,  0.])
```

```
# Our vectorized labels  
y_train = np.asarray(train_labels).astype('float32')  
y_test = np.asarray(test_labels).astype('float32')
```

Activation Function

```
from keras import models
from keras import layers
model = models.Sequential()
model.add(layers.Dense(16, activation='relu',
                       input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

$$\sum_{k=1}^{16} z_{1k} = \max(0, \sum_{j=1}^{10000} w_{1j}x_{1j} + b_1)$$

$$\sum_{l=1}^{16} z_{2l} = \max(0, \sum_{k=1}^{16} w_{2k}z_{1k} + b_2)$$

$$y = 1/\exp(-\sum_{l=1}^{16} w_{3l}z_{2l} + b_3)$$

Optimizer

```
from keras import optimizers  
  
model.compile(optimizer=optimizers.RMSprop(lr=0.001),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

$$r \leftarrow \rho r + (1 - \rho)g \odot g$$

$$\Delta W = -\frac{\epsilon}{\sqrt{\delta + r}} \odot g$$

$$-\hat{p}_0 \log(\hat{p}_0) - \hat{p}_1 \log(\hat{p}_1)$$

$$\hat{p}_c = \frac{1}{N} \sum_{n=1}^N I(y_i = c), \quad c = 0, 1$$

Validating the Approach

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
```

```
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

Train on 15000 samples, validate on 10000 samples

```
Epoch 1/20
15000/15000 [=====] - 1s - loss: 0.5103 - acc: 0.7911 - val_loss: 0.4016 - val_acc: 0.8628
Epoch 2/20
15000/15000 [=====] - 1s - loss: 0.3110 - acc: 0.9031 - val_loss: 0.3085 - val_acc: 0.8870
Epoch 3/20
15000/15000 [=====] - 1s - loss: 0.2309 - acc: 0.9235 - val_loss: 0.2803 - val_acc: 0.8908
Epoch 4/20
15000/15000 [=====] - 1s - loss: 0.1795 - acc: 0.9428 - val_loss: 0.2735 - val_acc: 0.8893
Epoch 5/20
```


History Object

```
history_dict = history.history  
history_dict.keys()
```

```
dict_keys(['val_acc', 'acc', 'val_loss', 'loss'])
```

```
import matplotlib.pyplot as plt  
  
acc = history.history['acc']  
val_acc = history.history['val_acc']  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs = range(1, len(acc) + 1)  
  
# "bo" is for "blue dot"  
plt.plot(epochs, loss, 'bo', label='Training loss')  
# b is for "solid blue line"  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()
```

Training and Validation Loss




Training and Validation Accuracy

```
plt.clf() # clear figure
acc_values = history_dict['acc']
val_acc_values = history_dict['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Evaluate on Test Data

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

Epoch 1/4

25000/25000 [=====] - 1s - loss: 0.4738 - acc: 0.8044

Epoch 2/4

25000/25000 [=====] - 1s - loss: 0.2660 - acc: 0.9076

Epoch 3/4

25000/25000 [=====] - 1s - loss: 0.2028 - acc: 0.9277

Epoch 4/4

25000/25000 [=====] - 1s - loss: 0.1700 - acc: 0.9397

24544/25000 [=====>.] - ETA: 0s

Generate Predictions

```
model.predict(x_test)
```

```
array([[ 0.91966152],  
       [ 0.86563045],  
       [ 0.99936908],  
       ...,  
       [ 0.45731062],  
       [ 0.0038014 ],  
       [ 0.79525089]], dtype=float32)
```